

Version: 1.1



Inception

Summary

This document is a System Administration-related exercise.

#Docker

#Administration

#System

42

Intellectual Property Disclaimer

All content presented in this training module, including but not limited to texts, images, graphics, and other materials, is protected by intellectual property rights held by Association 42.

Terms of Use:

- **Personal use:** You are permitted to use the contents of this module solely for personal purpose. Any commercial use, reproduction, distribution, modification, or public display is strictly prohibited without prior written permission from Association 42.
- **Respect for Integrity:** You must not alter, transform, or adapt the content in any way that could harm its integrity.

Protection of Rights:

Any violation of these terms constitutes an infringement of intellectual property rights and may result in legal action. We reserve the right to take all necessary measures to protect our rights, including but not limited to claims for damages.

For any questions regarding the use of the content or to obtain authorization, please contact:
legal@42.fr

Contents

1	Preamble	1
2	Introduction	2
3	AI Instructions	3
4	General guidelines	5
5	Mandatory part	6
6	Readme Requirements	11
7	Prerequisites for validation	12
8	Submission and peer-review	13

Chapter 1

Preamble



Chapter 2

Introduction

This activity aims to broaden your knowledge of system administration by using Docker. You will virtualize several Docker images, creating them in your new personal virtual machine.

Chapter 3

AI Instructions

● Context

During your learning journey, AI can assist with many different tasks. Take the time to explore the various capabilities of AI tools and how they can support your work. However, always approach them with caution and critically assess the results. Whether it's code, documentation, ideas, or technical explanations, you can never be completely sure that your question was well-formed or that the generated content is accurate. Your peers are a valuable resource to help you avoid mistakes and blind spots.

● Main message

- 👉 Use AI to reduce repetitive or tedious tasks.
- 👉 Develop prompting skills — both coding and non-coding — that will benefit your future career.
- 👉 Learn how AI systems work to better anticipate and avoid common risks, biases, and ethical issues.
- 👉 Continue building both technical and power skills by working with your peers.
- 👉 Only use AI-generated content that you fully understand and can take responsibility for.

● Learner rules:

- You should take the time to explore AI tools and understand how they work, so you can use them ethically and reduce potential biases.
- You should reflect on your problem before prompting — this helps you write clearer, more detailed, and more relevant prompts using accurate vocabulary.
- You should develop the habit of systematically checking, reviewing, questioning, and testing anything generated by AI.
- You should always seek peer review — don't rely solely on your own validation.

● Phase outcomes:

- Develop both general-purpose and domain-specific prompting skills.
- Boost your productivity with effective use of AI tools.
- Continue strengthening computational thinking, problem-solving, adaptability, and collaboration.

● Comments and examples:

- You'll regularly encounter situations — exams, evaluations, and more — where you must demonstrate real understanding. Be prepared, keep building both your technical and interpersonal skills.
- Explaining your reasoning and debating with peers often reveals gaps in your understanding. Make peer learning a priority.
- AI tools often lack your specific context and tend to provide generic responses. Your peers, who share your environment, can offer more relevant and accurate insights.
- Where AI tends to generate the most likely answer, your peers can provide alternative perspectives and valuable nuance. Rely on them as a quality checkpoint.

✓ Good practice:

I ask AI: "How do I test a sorting function?" It gives me a few ideas. I try them out and review the results with a peer. We refine the approach together.

✗ Bad practice:

I ask AI to write a whole function, copy-paste it into my activity. During peer-evaluation, I can't explain what it does or why. I lose credibility — and I fail my activity.

✓ Good practice:

I use AI to help design a parser. Then I walk through the logic with a peer. We catch two bugs and rewrite it together — better, cleaner, and fully understood.

✗ Bad practice:

I let Copilot generate my code for a key part of my activity. It compiles, but I can't explain how it handles pipes. During the evaluation, I fail to justify and I fail my activity.

Chapter 4

General guidelines

- This activity needs to be done on a Virtual Machine.
- All the files required for the configuration of your project must be placed in a `srcs` folder.
- A `Makefile` is also required and must be located at the root of your directory. It must set up your entire application (i.e., it has to build the Docker images using `docker-compose.yml`).
- This subject requires putting into practice concepts that, depending on your background, you may not have learned yet. Therefore, we advise you not to hesitate to read a lot of documentation related to Docker usage, as well as anything else you will find helpful in order to complete this assignment.

Chapter 5

Mandatory part

This activity consists of having you set up a small infrastructure composed of different services under specific rules. The whole project has to be done in a virtual machine. You have to use `docker compose`.

Each Docker image must have the same name as its corresponding service.

Each service has to run in a dedicated container.

For performance reasons, the containers must be built either from the penultimate stable version of Alpine or Debian. The choice is yours.

You also have to write your own Dockerfiles, one per service. The Dockerfiles must be called in your `docker-compose.yml` by your Makefile.

This means you have to build the Docker images of your project yourself. It is then forbidden to pull ready-made Docker images, as well as using services such as DockerHub (Alpine/Debian being excluded from this rule).

You then have to set up:

- A Docker container that contains NGINX with TLSv1.2 or TLSv1.3 only.
- A Docker container that contains WordPress + php-fpm (it must be installed and configured) only, without nginx.
- A Docker container that contains MariaDB only, without nginx.
- A volume that contains your WordPress database.
- A second volume that contains your WordPress website files.
- A `docker-network` that establishes the connection between your containers.

Your containers have to restart in case of a crash.



A Docker container is not a virtual machine. Thus, it is not recommended to use any hacky patch based on 'tail -f' and so forth when trying to run it. Read about how daemons work and whether it's a good idea to use them or not.



Of course, using `network: host` or `--link` or `links:` is forbidden. The `network` line must be present in your `docker-compose.yml` file. Your containers mustn't be started with a command running an infinite loop. Thus, this also applies to any command used as `entrypoint`, or used in `entrypoint` scripts. The following are a few prohibited hacky patches: `tail -f`, `bash`, `sleep infinity`, `while true`.



Read about PID 1 and the best practices for writing Dockerfiles.

- In your WordPress database, there must be two users, one of them being the administrator. The administrator's username can't contain `admin/Admin` or `administrator/Administrator` (e.g., `admin`, `administrator`, `Administrator`, `admin-123`, and so forth).



Your volumes will be available in the `/home/login/data` folder of the host machine using Docker. Of course, you have to replace the `login` with yours.

To make things simpler, you have to configure your domain name so it points to your local IP address.

This domain name must be `login.42.fr`. Again, you have to use your own `login`.

For example, if your `login` is `wil`, `wil.42.fr` will redirect to the IP address pointing to `wil`'s website.



The `latest` tag is prohibited.


No password must be present in your Dockerfiles.

It is mandatory to use environment variables.

Also, it is mandatory to use a `.env` file to store environment variables. It is strongly recommended that you use Docker `secrets` to store any confidential information. Any credentials, API keys, or passwords found in your Git repository (outside of properly configured secrets) will result in project failure.

Your NGINX container must be the only `entrypoint` into your infrastructure via the port 443 only, using the TLSv1.2 or TLSv1.3 protocol.

Here is an example diagram of the expected result:



Below is an example of the expected directory structure:

listing_part1

```
$> ls -alR
total XX
drwxrwxr-x  3 wil wil      4096 avril 42 20:42 .
drwxrwxrwt 17 wil wil      4096 avril 42 20:42 ..
-rw-rw-r-   1 wil wil      XXXX avril 42 20:42 Makefile
drwxrwxr-x  3 wil wil      4096 avril 42 20:42 secrets
drwxrwxr-x  3 wil wil      4096 avril 42 20:42 srcs
./secrets:
total XX
drwxrwxr-x  2 wil wil      4096 avril 42 20:42 .
drwxrwxr-x  6 wil wil      4096 avril 42 20:42 ..
-rw-r-r-   1 wil wil      XXXX avril 42 20:42 credentials.txt
-rw-r-r-   1 wil wil      XXXX avril 42 20:42 db_password.txt
-rw-r-r-   1 wil wil      XXXX avril 42 20:42 db_root_password.txt
./srcs:
total XX
drwxrwxr-x  3 wil wil      4096 avril 42 20:42 .
drwxrwxr-x  3 wil wil      4096 avril 42 20:42 ..
-rw-rw-r-   1 wil wil      XXXX avril 42 20:42 docker-compose.yml
-rw-rw-r-   1 wil wil      XXXX avril 42 20:42 .env
drwxrwxr-x  5 wil wil      4096 avril 42 20:42 requirements
./srcs/requirements:
total XX
drwxrwxr-x  5 wil wil      4096 avril 42 20:42 .
drwxrwxr-x  3 wil wil      4096 avril 42 20:42 ..
drwxrwxr-x  4 wil wil      4096 avril 42 20:42 bonus
drwxrwxr-x  4 wil wil      4096 avril 42 20:42 mariadb
drwxrwxr-x  4 wil wil      4096 avril 42 20:42 nginx
drwxrwxr-x  4 wil wil      4096 avril 42 20:42 tools
drwxrwxr-x  4 wil wil      4096 avril 42 20:42 wordpress
```

listing_part2

```
./srcs/requirements/mariadb:
total XX
drwxrwxr-x  4 wil wil      4096 avril 42 20:45 .
drwxrwxr-x  5 wil wil      4096 avril 42 20:42 ..
drwxrwxr-x  2 wil wil      4096 avril 42 20:42 conf
-rw-rw-r-   1 wil wil      XXXX avril 42 20:42 Dockerfile
-rw-rw-r-   1 wil wil      XXXX avril 42 20:42 .dockerignore
drwxrwxr-x  2 wil wil      4096 avril 42 20:42 tools
[...]
./srcs/requirements/nginx:
total XX
drwxrwxr-x  4 wil wil      4096 avril 42 20:42 .
drwxrwxr-x  5 wil wil      4096 avril 42 20:42 ..
drwxrwxr-x  2 wil wil      4096 avril 42 20:42 conf
-rw-rw-r-   1 wil wil      XXXX avril 42 20:42 Dockerfile
-rw-rw-r-   1 wil wil      XXXX avril 42 20:42 .dockerignore
drwxrwxr-x  2 wil wil      4096 avril 42 20:42 tools
[...]
$> cat srcs/.env
DOMAIN_NAME=wil.42.fr
# MYSQL SETUP
MYSQL_USER=XXXXXXXXXXXXX
[...]
$>
```



For obvious security reasons, any credentials, API keys, passwords, etc., must be saved locally in various ways/files and ignored by git. Publicly stored credentials will lead you directly to a failure of the activity.



You can store your variables (as a domain name) in an environment variable file like .env

Chapter 6

Readme Requirements

A README.md file must be provided at the root of your Git repository. Its purpose is to allow anyone unfamiliar with the activity (peers, staff, recruiters, etc.) to quickly understand what the activity is about, how to run it, and where to find more information on the topic.

The README.md must include at least:

- The very first line must be italicized and read: *This activity has been created as part of the 42 curriculum by <login1>[, <login2>[, <login3>[...]]]*.
- A "**Description**" section that clearly presents the activity, including its goal and a brief overview.
- An "**Instructions**" section containing any relevant information about compilation, installation, and/or execution.
- A "**Resources**" section listing classic references related to the topic (documentation, articles, tutorials, etc.), as well as a description of how AI was used — specifying for which tasks and which parts of the activity.

➡ **Additional sections may be required depending on the activity** (e.g., usage examples, feature list, technical choices, etc.).

Any required additions will be explicitly listed below.

- A **Project description** section must also explain the use of Docker and the sources included in the activity. It must indicate the main design choices, as well as a comparison between:
 - Virtual Machines vs Docker
 - Secrets vs Environment Variables
 - Docker Network vs Host Network
 - Docker Volumes vs Bind Mounts



Your README must be written in English.

Chapter 7

Prerequisites for validation

In addition to the existing requirements, the following documentation files must be present at the root of your repository. They must be written in Markdown format (.md).

- **USER_DOC.md — User documentation** This file must explain, in clear and simple terms, how an end user or administrator can:
 - Understand what services are provided by the stack.
 - Start and stop the activity.
 - Access the website and the administration panel.
 - Locate and manage credentials.
 - Check that the services are running correctly.
- **DEV_DOC.md — Developer documentation** This file must describe how a developer can:
 - Set up the environment from scratch (prerequisites, configuration files, secrets).
 - Build and launch the activity using the Makefile and Docker Compose.
 - Use relevant commands to manage the containers and volumes.
 - Identify where the activity data is stored and how it persists.

Chapter 8

Submission and peer-review

Submit your assignment in your `Git` repository as usual. Only the work inside your repository will be evaluated during the defense. Don't hesitate to double-check the names of your folders and files to ensure they are correct.

During the review, a brief **modification of the activity** may occasionally be requested. This could involve a minor behavior change, a few lines of code to write or rewrite, or an easy-to-add feature.

While this step may **not be applicable to every activity**, you must be prepared for it if it is mentioned in the review guidelines.

This step is meant to verify your actual understanding of a specific part of the activity. The modification can be performed in any development environment you choose (e.g., your usual setup), and it should be feasible within a few minutes — unless a specific timeframe is defined as part of the review.

You can, for example, be asked to make a small update to a function or script, modify a display, or adjust a data structure to store new information, etc.

The details (scope, target, etc.) will be specified in the **review guidelines** and may vary from one review to another for the same activity.